

HMF Design Clinic Grading Guide

The rubric proctors use to grade each team's submitted Workbook

This is a reference guide, not a worksheet. It defines, for every answer, the levels of quality mapped to a score, so any proctor grades consistently. Read a team's answer, classify it against the rubric, and take the score for that level. Record the scores on the Proctor evaluation page at the back of the team's Workbook. Use the Solution Guide for the wider model answers.

How the score is built

Every pain point is graded on the three dimensions. Problem Identifier is worth 5, Solution Designer 10, and Value Communicator 10, for 25 per pain point. Four pain points make a case study score out of 100, and the four case studies make an overall score out of 400.

Problem Identifier OUT

OF 5

Correctly identified the customer's pain point and the requirements behind it.

Solution Designer OUT

OF 10

Recommended the right Cisco solution, architecture, and approach.

Value Communicator

OUT OF 10

Clearly explained how the solution works and the customer outcomes it enables.

How to use this guide. For each answer, decide which level in the rubric it matches on each dimension, from Complete down to Minimal, and take the points for that level. Add the three dimensions for the pain point total, and the four pain points for the case study score. Grade the reasoning, not the wording: a different but well-argued product can still earn marks if it fits Cedarline. Record every score on the team's Proctor evaluation page.

Levels of an answer, by dimension

Each dimension has four levels. Match the team's written answer to the level it best fits, and award points from that level's range. These levels apply to every pain point; the full-marks answer for each specific pain point is given in the answer key that follows.

Problem Identifier OUT OF 5

Did the team identify the real requirement behind the concern?

LEVEL	POINTS	WHAT THE TEAM'S ANSWER LOOKS LIKE AT THIS LEVEL
Complete	5	Pinpoints the underlying requirement or constraint behind the concern, the specific regulation, risk, or operational need, in Cedarline's own context, not just the surface ask.
Substantial	3 to 4	Identifies the concern and part of the underlying requirement, but misses some nuance or a constraint.
Developing	1 to 2	Restates the concern in general terms without naming the requirement behind it.
Minimal	0	Misreads the concern, addresses the wrong problem, or leaves it blank.

Solution Designer OUT OF 10

Did the team recommend the right Cisco solution and an approach that fits?

LEVEL	POINTS	WHAT THE TEAM'S ANSWER LOOKS LIKE AT THIS LEVEL
Complete	8 to 10	Names the right Cisco product or platform and the specific capability, with an architecture or approach that fits Cedarline and the wider Hybrid Mesh Firewall design.
Substantial	5 to 7	Names the right product but is light on the capability, the architecture, or how it fits Cedarline.
Developing	2 to 4	Partly right: a loosely related product, or a product named with no approach or capability.
Minimal	0 to 1	Wrong product, a generic firewall with no specifics, or left blank.

Value Communicator OUT OF 10

Did the team explain how it works and the outcome for Cedarline?

LEVEL	POINTS	WHAT THE TEAM'S ANSWER LOOKS LIKE AT THIS LEVEL
Complete	8 to 10	Explains clearly how the solution works and states the concrete Cedarline outcome, in language the customer would understand.
Substantial	5 to 7	Explains how it works but is thin on the outcome, or states an outcome without explaining how.

LEVEL	POINTS	WHAT THE TEAM'S ANSWER LOOKS LIKE AT THIS LEVEL
Developing	2 to 4	Vague or generic: asserts it helps without explaining how or naming a specific outcome.
Minimal	0 to 1	No explanation and no outcome, or left blank.

From completeness to a score

Use these bands to sense-check the totals. A pain point total should reflect how complete the answer is across the three dimensions, and the case study score should reflect the four pain points together.

Pain point, out of 25

LEVEL	OUT OF 25	WHAT IT MEANS
Complete	22 to 25	All three dimensions strong: the requirement, the right solution and approach, and a clear how and outcome.
Substantial	16 to 21	Mostly there, with one dimension light or a gap in the approach or the outcome.
Developing	9 to 15	Partial across the dimensions: the idea is present but thin, generic, or missing the requirement or the solution.
Minimal	0 to 8	Little relevant content, or the main dimensions left blank.

Case study, out of 100

LEVEL	OUT OF 100	WHAT IT MEANS
Exemplary	85 to 100	Consistently complete answers across the four pain points, specific to Cedarline.
Proficient	65 to 84	Solid answers, with some gaps or generic reasoning in places.
Developing	40 to 64	Partial: the direction is right, but several answers are thin or incomplete.
Emerging	under 40	Mostly incomplete or off-target answers.

Recording the score. This guide is for grading only; it has no spaces to write in. Record each case study score, the overall score out of 400, and your feedback on the **Proctor evaluation page** at the back of the team's Workbook.

Working through a Workbook

Each team submits a completed Workbook. Work through it one case study at a time, grading every pain point against the rubric and the answer key on the pages that follow.

- 1 Collect each team's Workbook and check that all four case studies have been attempted.
- 2 Take one case study at a time. For each pain point, read the team's Product / platform and Solution Engineer response.
- 3 For each dimension, classify the team's answer into a level using the rubric: Complete, Substantial, Developing or Minimal, comparing against the full-credit answer for that pain point.
- 4 Take the points for the level on each dimension (Problem Identifier out of 5, Solution Designer out of 10, Value Communicator out of 10) and add them for the pain point total out of 25.
- 5 Add the four pain point totals for the case study score out of 100, and read the overall completion band to sense-check the grade.
- 6 Record each case study score and a line of strengths and coaching on the Proctor evaluation page at the back of the team's Workbook.
- 7 Add the four case study scores for the overall score out of 400, note what the team did well and where they can grow, and pass the topology diagram to the SME for review.

What each dimension looks for. Problem Identifier: did the team name the real requirement behind the concern? **Solution Designer:** did they land on the right product and an approach that fits Cedarline? **Value Communicator:** did they explain how it works and the outcome? Compare against the answer key, and reward reasoning specific to Cedarline.

1

Case Study 1

The answer key below is the **Complete** (full-marks) answer for each pain point. Grade the team's answer against it and the rubric levels; a Complete answer earns the top of each dimension's range.

PAIN POINT 1 Answer key, full marks 25

Under the EU AI Act, we are classified as a 'Provider' for our internal AI systems. We need to demonstrate robustness and cybersecurity compliance to satisfy Article 15. How does Cisco help us meet these specific requirements?

Problem Identifier OUT OF 5	Provider-level evidence of robustness and cybersecurity under EU AI Act Article 15.
Solution Designer OUT OF 10	Cisco AI Defense (AI Model and Application Validation)
Value Communicator OUT OF 10	How it works. Cisco AI Defense is specifically designed to support Provider obligations under the EU AI Act. We utilize algorithmic red teaming technology for AI Model & App Validation, which identifies safety and security vulnerabilities at scale. Outcome. It automates the technical evidence gathering required for "Conformity Assessments." By mapping AI-specific threats to standardized frameworks (like MITRE ATLAS and OWASP), it provides the documentation and continuous monitoring necessary to prove to regulators that your systems are robust and compliant.

PAIN POINT 2 Answer key, full marks 25

We are integrating various open-source models and third-party agents into our internal workflows. How can we ensure these assets aren't introducing vulnerabilities into our environment before they even go live?

Problem Identifier OUT OF 5	Vetting AI models and third-party agents for vulnerabilities before they reach production.
Solution Designer OUT OF 10	Cisco AI Defense (Supply Chain Risk Management)
Value Communicator OUT OF 10	How it works. Cisco AI Defense addresses this through proactive Supply Chain Risk Management. We scan your model files, repositories, and MCP servers to identify malicious or unsafe assets. Outcome. It enables "Shift-Left" security by creating an inventory of all AI models and agents. By identifying vulnerabilities before production, it prevents the introduction of malicious code or poisoned data into your ecosystem.

PAIN POINT 3

Answer key, full marks 25

We know AI is showing up across our cloud environments faster than security can track it. We're worried we have models, agents, datasets, and MCP-connected tools running without visibility or consistent controls.

Problem Identifier OUT OF 5	Visibility and consistent control over AI assets (models, agents, datasets and MCP tools) across the cloud.
Solution Designer OUT OF 10	Cisco AI Defense (AI Cloud Visibility)
Value Communicator OUT OF 10	How it works. That's exactly the visibility gap Cisco AI Defense is built to solve. With AI Cloud Visibility, Cisco AI Defense continuously discovers AI assets across distributed cloud environments, including models, agents, datasets, MCP servers, and agent-to-tool workflows. Outcome. Identifies rogue or unsanctioned AI assets. Maps relationships across data, models, agents, and tools.

PAIN POINT 4

Answer key, full marks 25

We're running AI workloads in cloud accounts, but we don't have a reliable way to see which models and AI assets are actually in use. We're also concerned about workloads reaching out to third-party AI models without security knowing about it. And even when we do know the traffic path, we need a practical way to inspect and enforce guardrails on prompt and response traffic in the cloud without changing every application.

Problem Identifier OUT OF 5	Seeing AI assets and third-party model usage, and enforcing runtime guardrails on AI traffic without changing applications.
Solution Designer OUT OF 10	Cisco AI Defense with Cisco Multicloud Defense (AI Runtime Protection)
Value Communicator OUT OF 10	How it works. That's exactly what the Cisco AI Defense-Cisco Multicloud Defense integration is designed to address. Cisco AI Defense uses the integration to discover AI assets in your cloud environments and identify external AI destinations your workloads are calling. Outcome. Finds AI assets in cloud accounts. Identifies third-party AI model usage.

2

Case Study 2

The answer key below is the **Complete** (full-marks) answer for each pain point. Grade the team's answer against it and the rubric levels; a Complete answer earns the top of each dimension's range.

PAIN POINT 1 Answer key, full marks 25

If a breach occurs in one of our public cloud instances, we are terrified that the attacker will move laterally into our core on-premises systems. How can we contain these threats in a multicloud environment?

Problem Identifier OUT OF 5	Containing lateral movement between public cloud and on-prem in a multicloud environment.
Solution Designer OUT OF 10	Cisco Multicloud Defense (macro and micro-segmentation gateways)
Value Communicator OUT OF 10	How it works. Cisco Multicloud Defense leverages macro and micro-segmentation capabilities to enforce strict communication boundaries. By integrating with our Hybrid Mesh Firewall and Hypershield, we can isolate workloads and inspect traffic between cloud segments, effectively stopping lateral movement before it reaches your sensitive data. Outcome. It provides granular visibility and control over east-west traffic. By applying Zero Trust principles across the hybrid fabric, you transform your network into a series of secure, isolated zones that prevent an attacker from pivoting through your infrastructure.

PAIN POINT 2 Answer key, full marks 25

Our auditors require proof that our security controls are applied consistently across all our cloud workloads. Gathering this data from different cloud providers is a manual, time-consuming nightmare.

Problem Identifier OUT OF 5	Consistent, auditable proof that security controls are applied across all cloud workloads.
Solution Designer OUT OF 10	Cisco Multicloud Defense (unified logging and visibility)
Value Communicator OUT OF 10	How it works. Cisco Multicloud Defense provides unified logging and traceability. Because our security framework is integrated across the network fabric, you can generate consolidated reports that map directly to your compliance requirements (such as NIS2 or DORA) across all cloud environments from a single dashboard. Outcome. It automates the collection of evidence for security audits. By consolidating logs and policy enforcement data, you create a "single source of truth" that simplifies compliance reporting and demonstrates proactive governance to regulators.

PAIN POINT 3

Answer key, full marks 25

We have workloads spread across AWS, Azure, and on-premises data centers. Managing security policies individually in each cloud environment is creating massive operational overhead and leading to inconsistent security posture.

Problem Identifier OUT OF 5	One consistent policy model across AWS, Azure and on-prem, without per-cloud overhead.
Solution Designer OUT OF 10	Cisco Multicloud Defense (unified policy across public clouds)
Value Communicator OUT OF 10	How it works. Cisco Multicloud Defense, integrated within our Hybrid Mesh Firewall architecture, provides a unified control plane. You can define security policies once via Security Cloud Control and enforce them consistently across your entire distributed environment, regardless of the underlying cloud provider. Outcome. It replaces siloed, manual configuration with centralized orchestration. By abstracting the security policy from the cloud-native tools, it ensures that your security posture is uniform and eliminates the "configuration drift" that occurs when managing multiple cloud consoles.

PAIN POINT 4

Answer key, full marks 25

We have workloads on-prem, in cloud, and in mixed environments, and policy consistency is becoming difficult. Every environment feels different, and we're worried about drift.

Problem Identifier OUT OF 5	Unified segmentation that prevents policy drift across hybrid and multicloud.
Solution Designer OUT OF 10	Cisco Secure Workload with Cisco Secure Firewall (managed via FMC and Security Cloud Control)
Value Communicator OUT OF 10	How it works. That's a common hybrid environment challenge. Cisco Secure Workload helps create a more unified segmentation model by using application dependency awareness and AI/ML-driven policy discovery to define policy based on how workloads actually communicate. Outcome. Reduces policy drift across hybrid and multicloud environments. Aligns policy to application behavior, not just static IP-based rules.

3 Case Study 3

The answer key below is the **Complete** (full-marks) answer for each pain point. Grade the team's answer against it and the rubric levels; a Complete answer earns the top of each dimension's range.

PAIN POINT 1 Answer key, full marks 25

When our security tools detect a compromised device, our manual response time is too slow. By the time we identify the user and manually isolate the port, the attacker has already moved through the network.

Problem Identifier OUT OF 5	Automated, real-time isolation of a compromised device, faster than a manual response.
Solution Designer OUT OF 10	Cisco ISE (Security Group Tags via pxGrid and Adaptive Network Control)
Value Communicator OUT OF 10	How it works. Cisco ISE integrates with your security ecosystem via pxGrid and Adaptive Network Control (ANC). When a threat is detected, ISE can automatically trigger a quarantine policy, isolating the compromised device in real-time without requiring manual intervention. Outcome. It automates the "Response" phase of the incident lifecycle. By bridging the gap between detection and enforcement, it stops threats at the speed of the network.

PAIN POINT 2 Answer key, full marks 25

Our firewall environment has become too complex to manage manually.

Problem Identifier OUT OF 5	Simplifying and cleaning up firewall policy administration at scale.
Solution Designer OUT OF 10	Cisco Security Cloud Control with FMC and Cisco Secure Firewall (unified policy administration)
Value Communicator OUT OF 10	How it works. Cisco applies AI-driven analysis to firewall telemetry and policy data to surface optimization opportunities, reduce manual troubleshooting, and guide cleaner policy administration. Outcome. Reduces manual policy review; identifies redundant/stale rules; lowers risk from human error.

Most of our traffic is encrypted, so we've lost visibility into a huge part of our environment.

Problem Identifier OUT OF 5	Restoring visibility into encrypted traffic without decrypting everything.
Solution Designer OUT OF 10	Cisco Secure Firewall with the Encrypted Visibility Engine (EVE)
Value Communicator OUT OF 10	How it works. Encrypted Visibility Engine (EVE) uses AI/ML-based fingerprinting to analyze encrypted traffic patterns and identify suspicious behavior without decrypting the payload. Outcome. Restores visibility into encrypted blind spots; avoids the performance/privacy overhead of decrypt-everything strategies.

We're worried traditional signature-based detection won't catch new or unknown attacks quickly enough.

Problem Identifier OUT OF 5	Detecting unknown and zero-day threats beyond signature-based detection.
Solution Designer OUT OF 10	Cisco Secure Firewall with Snort ML
Value Communicator OUT OF 10	How it works. That's exactly where Snort ML helps. It adds machine-learning-based detection inside the Snort 3 engine to identify exploit behavior associated with unknown/zero-day threats. Outcome. Addresses the gap between known and unknown threat detection; reduces dependence on signature timing.

4

Case Study 4

The answer key below is the **Complete** (full-marks) answer for each pain point. Grade the team's answer against it and the rubric levels; a Complete answer earns the top of each dimension's range.

PAIN POINT 1 Answer key, full marks 25

Network policy helps, but we also need to know what's happening at runtime inside the workload. We need better detection of suspicious behavior and policy violations.

Problem Identifier OUT OF 5	Runtime detection of suspicious behaviour and policy violations inside workloads.
Solution Designer OUT OF 10	Isovalent Tetragon (eBPF runtime security)
Value Communicator OUT OF 10	How it works. That's where Tetragon comes in. Tetragon extends the value of Cilium by providing eBPF-based runtime security and threat detection. Outcome. Adds runtime visibility beyond basic network controls. Detects suspicious system and network behavior in real time.

PAIN POINT 2 Answer key, full marks 25

Our Kubernetes networking feels fragmented across clusters and environments. We're dealing with too many moving parts, inconsistent policies, and too much operational complexity.

Problem Identifier OUT OF 5	Consistent Kubernetes networking and policy across clusters, with less operational complexity.
Solution Designer OUT OF 10	Isovalent Cilium (eBPF Kubernetes networking and policy)
Value Communicator OUT OF 10	How it works. That's exactly where Isovalent Cilium helps. Cilium provides a modern, eBPF-based Kubernetes networking layer that gives you consistent connectivity, policy enforcement, and visibility across clusters and environments. Outcome. Replaces brittle, IP-centric networking models with a Kubernetes-native approach. Improves consistency across dynamic cluster environments.

PAIN POINT 3

Answer key, full marks 25

Right now we're stitching together separate tools for networking, security, and observability in Kubernetes, and it's creating too much overhead.

Problem Identifier OUT OF 5	Consolidating Kubernetes networking, security and observability to reduce tool sprawl.
Solution Designer OUT OF 10	Isovalent Cilium with Tetragon
Value Communicator OUT OF 10	How it works. That's exactly the consolidation opportunity Isovalent addresses. With Cilium for networking and policy, and Tetragon for runtime security, you get a more unified platform for connectivity, segmentation, observability, and security. Outcome. Consolidates multiple functions into one platform approach. Reduces tool sprawl and operational friction.

PAIN POINT 4

Answer key, full marks 25

We know we need microsegmentation, but we don't have enough visibility into east-west flows or application dependencies to create policy safely. We don't want to break the application.

Problem Identifier OUT OF 5	Mapping east-west flows and application dependencies to segment safely without breaking applications.
Solution Designer OUT OF 10	Cisco Secure Workload (application dependency mapping)
Value Communicator OUT OF 10	How it works. That's one of the biggest reasons customers adopt Cisco Secure Workload. It gives you deep visibility into workload communication and application dependency mapping, so you can see how applications actually behave before enforcing policy. Outcome. Reveals actual application communication paths. Reduces the risk of segmentation causing outages.