

HMF Design Clinic Grading Guide

The rubric proctors use to grade each team's submitted Workbook

This is a reference guide, not a worksheet. It defines, for every answer, the levels of quality mapped to a score, so any proctor grades consistently. Read a team's answer, classify it against the rubric, and take the score for that level. Record the scores on the Proctor evaluation page at the back of the team's Workbook. Use the Solution Guide for the wider model answers.

How the score is built

Every pain point is graded on the three dimensions. Problem Identifier is worth 5, Solution Designer 10, and Value Communicator 10, for 25 per pain point. Four pain points make a case study score out of 100, and the four case studies make a cumulative score out of 400.

Problem Identifier OUT
OF 5

Correctly identified the customer's pain point and the requirements behind it.

Solution Designer OUT
OF 10

Recommended the right Cisco solution, architecture, and approach.

Value Communicator
OUT OF 10

Clearly explained how the solution works and the customer outcomes it enables.

How to use this guide. For each answer, decide which level in the rubric it matches on each dimension, from Complete down to Minimal, and take the points for that level. Add the three dimensions for the pain point total, and the four pain points for the case study score. Grade the reasoning, not the wording: a different but well-argued product can still earn marks if it fits Cedarline. Record every score on the team's Proctor evaluation page.

Levels of an answer, by dimension

Each dimension has four levels. Match the team's written answer to the level it best fits, and award points from that level's range. These levels apply to every pain point; the full-marks answer for each specific pain point is given in the answer key that follows.

Problem Identifier OUT OF 5

Did the team identify the real requirement behind the concern?

LEVEL	POINTS	WHAT THE TEAM'S ANSWER LOOKS LIKE AT THIS LEVEL
Complete	5	Pinpoints the underlying requirement or constraint behind the concern, the specific regulation, risk, or operational need, in Cedarline's own context, not just the surface ask.
Substantial	3 to 4	Identifies the concern and part of the underlying requirement, but misses some nuance or a constraint.
Developing	1 to 2	Restates the concern in general terms without naming the requirement behind it.
Minimal	0	Misreads the concern, addresses the wrong problem, or leaves it blank.

Solution Designer OUT OF 10

Did the team recommend the right Cisco solution and an approach that fits?

LEVEL	POINTS	WHAT THE TEAM'S ANSWER LOOKS LIKE AT THIS LEVEL
Complete	8 to 10	Names the right Cisco product or platform and the specific capability, with an architecture or approach that fits Cedarline and the wider Hybrid Mesh Firewall design.
Substantial	5 to 7	Names the right product but is light on the capability, the architecture, or how it fits Cedarline.
Developing	2 to 4	Partly right: a loosely related product, or a product named with no approach or capability.
Minimal	0 to 1	Wrong product, a generic firewall with no specifics, or left blank.

Value Communicator OUT OF 10

Did the team explain how it works and the outcome for Cedarline?

LEVEL	POINTS	WHAT THE TEAM'S ANSWER LOOKS LIKE AT THIS LEVEL
Complete	8 to 10	Explains clearly how the solution works and states the concrete Cedarline outcome, in language the customer would understand.
Substantial	5 to 7	Explains how it works but is thin on the outcome, or states an outcome without explaining how.

LEVEL	POINTS	WHAT THE TEAM'S ANSWER LOOKS LIKE AT THIS LEVEL
Developing	2 to 4	Vague or generic: asserts it helps without explaining how or naming a specific outcome.
Minimal	0 to 1	No explanation and no outcome, or left blank.

From completeness to a score

Use these bands to sense-check the totals. A pain point total should reflect how complete the answer is across the three dimensions, and the case study score should reflect the four pain points together.

Pain point, out of 25

LEVEL	OUT OF 25	WHAT IT MEANS
Complete	22 to 25	All three dimensions strong: the requirement, the right solution and approach, and a clear how and outcome.
Substantial	16 to 21	Mostly there, with one dimension light or a gap in the approach or the outcome.
Developing	9 to 15	Partial across the dimensions: the idea is present but thin, generic, or missing the requirement or the solution.
Minimal	0 to 8	Little relevant content, or the main dimensions left blank.

Case study, out of 100

LEVEL	OUT OF 100	WHAT IT MEANS
Exemplary	85 to 100	Consistently complete answers across the four pain points, specific to Cedarline.
Proficient	65 to 84	Solid answers, with some gaps or generic reasoning in places.
Developing	40 to 64	Partial: the direction is right, but several answers are thin or incomplete.
Emerging	under 40	Mostly incomplete or off-target answers.

Overall grade, out of 400

Give the team their overall result as one of these three encouraging categories. Every category is positive; it names where the team is strong and what to build next.

LEVEL	OUT OF 400	WHAT IT MEANS
Advanced	300 to 400	Standout work. Across the four case studies the team read the requirements precisely, designed fitting Hybrid Mesh Firewall solutions, and told a clear Cedarline value story.
Established	200 to 299	Strong, consistent work. The team matched the right solutions across most pain points and communicated real value, with a few answers to sharpen.
Foundational	up to 199	A solid foundation to build on. The team engaged every case study and found real products; the next step is naming the requirement and the outcome more precisely.

Recording the score. This guide is for grading only; it has no spaces to write in. Record each case study score, the overall score out of 400 and its category, and your feedback on the **Proctor evaluation page** at the back of the team's Workbook.

Working through a Workbook

Each team submits a completed Workbook. Work through it one case study at a time, grading every pain point against the rubric and the answer key on the pages that follow.

- 1 Collect each team's Workbook and check that all four case studies have been attempted.
- 2 Take one case study at a time. For each pain point, read the team's Product / platform and Solution Engineer response.
- 3 For each dimension, classify the team's answer into a level using the rubric: Complete, Substantial, Developing or Minimal, comparing against the full-credit answer for that pain point.
- 4 Take the points for the level on each dimension (Problem Identifier out of 5, Solution Designer out of 10, Value Communicator out of 10) and add them for the pain point total out of 25.
- 5 Add the four pain point totals for the case study score out of 100, and read the overall completion band to sense-check the grade.
- 6 Record each case study score and a line of strengths and coaching on the Proctor evaluation page at the back of the team's Workbook.
- 7 Add the four case study scores for the overall score out of 400, note what the team did well and where they can grow, and pass the topology diagram to the SME for review.

What each dimension looks for. Problem Identifier: did the team name the real requirement behind the concern? **Solution Designer:** did they land on the right product and an approach that fits Cedarline? **Value Communicator:** did they explain how it works and the outcome? Compare against the answer key, and reward reasoning specific to Cedarline.

Grading alternative answers. The full-marks answer shows the primary Cisco fit, but a team may propose a different Hybrid Mesh Firewall component. Each pain point lists the **acceptable alternatives** under Solution Designer. Award full marks for any acceptable alternative when the reasoning fits Cedarline, partial when the product is related but the approach is thin, and no credit for a generic or wrong product. Grade sound reasoning over the exact product name.

Grading a sample answer, step by step

Before you grade real Workbooks, walk through this example. It shows how to move from a team's written answer to a score on one pain point, the same way you will grade all sixteen.

The pain point. *Under the EU AI Act, we are classified as a 'Provider' for our internal AI systems. We need to demonstrate robustness and cybersecurity compliance to satisfy Article 15. How does Cisco help us meet these specific requirements?*

THE TEAM WROTE

Product / platform. Cisco AI Defense

Solution Engineer response. Cisco AI Defense helps with AI security and can validate our models. It will help us stay compliant with regulations and keep our AI safe.

GRADE IT STEP BY STEP

- Step 1** **Grade the Problem Identifier answer.** It matches **Partial**, so score **3 out of 5**. The team points at compliance but does not name the Article 15 requirement to prove robustness and cybersecurity as a Provider. Right direction, missing the specific requirement, so it lands in the partial range.
- Step 2** **Grade the Solution Designer answer.** It matches **Substantial**, so score **6 out of 10**. Cisco AI Defense is the right product, but the team does not name the model and application validation capability or how it is applied. Right product, light on the approach.
- Step 3** **Grade the Value Communicator answer.** It matches **Substantial**, so score **6 out of 10**. They mention validating models, which is some of the how, but the outcome stays generic. It does not reach the concrete outcome, automated evidence for Conformity Assessments.
- Step 4** **Total the pain point.** 3 plus 6 plus 6 = **15 out of 25**, which is **Developing** on the completion band.
- Step 5** **Note the feedback.** Strength: the right product and a clear compliance intent. To grow: name the Article 15 requirement, the AI Defense model and application validation capability, and the concrete outcome, automated evidence for Conformity Assessments.

SAMPLE OUTPUT

Pain point 1: Problem 3 plus Solution 6 plus Value 6 = 15 out of 25, which is Developing. Grade pain points 2 to 4 the same way, add the four totals for the Case Study 1 score out of 100, and record it on the team's Proctor evaluation page.

Now you are ready. Grade each pain point this way against the three levels for every dimension, add the four pain point totals for the case study score out of 100, and record every score on the team's Proctor evaluation page.

1

Case Study 1

For each pain point, grade all three dimensions. Each shows the **full-marks** answer, what a **partial** answer looks like, and what earns **no credit**, mapped to the score.

PAIN POINT 1 Grade all three dimensions, full marks 25

Under the EU AI Act, we are classified as a 'Provider' for our internal AI systems. We need to demonstrate robustness and cybersecurity compliance to satisfy Article 15. How does Cisco help us meet these specific requirements?

Problem Identifier OUT OF 5

FULL, 5	Provider-level evidence of robustness and cybersecurity under EU AI Act Article 15.
PARTIAL, 1 TO 4	Sees the concern but names only part of it, or states it in general terms without the specific requirement above.
NO CREDIT, 0	Restates or misreads the concern; the underlying requirement is missing.

Solution Designer OUT OF 10

FULL, 8 TO 10	Cisco AI Defense (AI Model and Application Validation) Also acceptable: Cisco AI Defense (model validation and algorithmic red teaming) (the same platform named by capability; full marks when it is tied to proving robustness and cybersecurity for Article 15).
PARTIAL, 2 TO 7	Names Cisco AI Defense but does not connect it to Article 15 evidence or model validation.
NO CREDIT, 0 TO 1	A policy or governance process with no Cisco product, or an unrelated next-generation firewall.

Value Communicator OUT OF 10

FULL, 8 TO 10	How it works. Cisco AI Defense is specifically designed to support Provider obligations under the EU AI Act. We utilize algorithmic red teaming technology for AI Model & App Validation, which identifies safety and security vulnerabilities at scale. Outcome. It automates the technical evidence gathering required for "Conformity Assessments." By mapping AI-specific threats to standardized frameworks (like MITRE ATLAS and OWASP), it provides the documentation and continuous monitoring necessary to prove to regulators that your systems are robust and compliant.
PARTIAL, 2 TO 7	Explains the how or the outcome, but not both, or keeps it generic rather than specific to Cedarline.
NO CREDIT, 0 TO 1	Claims it helps, with no explanation of how it works and no Cedarline outcome.

We are integrating various open-source models and third-party agents into our internal workflows. How can we ensure these assets aren't introducing vulnerabilities into our environment before they even go live?

Problem Identifier OUT OF 5

FULL, 5	Vetting AI models and third-party agents for vulnerabilities before they reach production.
PARTIAL, 1 TO 4	Sees the concern but names only part of it, or states it in general terms without the specific requirement above.
NO CREDIT, 0	Restates or misreads the concern; the underlying requirement is missing.

Solution Designer OUT OF 10

FULL, 8 TO 10	Cisco AI Defense (Supply Chain Risk Management) Also acceptable: Cisco AI Defense applied as a pre-production gate (AI Defense validation run in the model and agent pipeline; full marks when the vet-before-production reasoning is sound).
PARTIAL, 2 TO 7	Names AI Defense but not the pre-production vetting of third-party models and agents.
NO CREDIT, 0 TO 1	Endpoint antivirus or code scanning with no AI supply-chain reasoning.

Value Communicator OUT OF 10

FULL, 8 TO 10	How it works. Cisco AI Defense addresses this through proactive Supply Chain Risk Management. We scan your model files, repositories, and MCP servers to identify malicious or unsafe assets. Outcome. It enables "Shift-Left" security by creating an inventory of all AI models and agents. By identifying vulnerabilities before production, it prevents the introduction of malicious code or poisoned data into your ecosystem.
PARTIAL, 2 TO 7	Explains the how or the outcome, but not both, or keeps it generic rather than specific to Cedarline.
NO CREDIT, 0 TO 1	Claims it helps, with no explanation of how it works and no Cedarline outcome.

We know AI is showing up across our cloud environments faster than security can track it. We're worried we have models, agents, datasets, and MCP-connected tools running without visibility or consistent controls.

Problem Identifier OUT OF 5

FULL, 5	Visibility and consistent control over AI assets (models, agents, datasets and MCP tools) across the cloud.
PARTIAL, 1 TO 4	Sees the concern but names only part of it, or states it in general terms without the specific requirement above.
NO CREDIT, 0	Restates or misreads the concern; the underlying requirement is missing.

Solution Designer OUT OF 10

FULL, 8 TO 10	Cisco AI Defense (AI Cloud Visibility) Also acceptable: Cisco Multicloud Defense (cloud traffic visibility) (complements AI Defense to discover AI assets and traffic across clouds; credit when paired or well argued).
PARTIAL, 2 TO 7	Names AI Defense but not the discovery of models, agents, datasets and MCP tools across the cloud.
NO CREDIT, 0 TO 1	A generic cloud logging or posture tool with no AI-asset discovery.

Value Communicator OUT OF 10

FULL, 8 TO 10	How it works. That's exactly the visibility gap Cisco AI Defense is built to solve. With AI Cloud Visibility, Cisco AI Defense continuously discovers AI assets across distributed cloud environments, including models, agents, datasets, MCP servers, and agent-to-tool workflows. Outcome. Identifies rogue or unsanctioned AI assets. Maps relationships across data, models, agents, and tools.
PARTIAL, 2 TO 7	Explains the how or the outcome, but not both, or keeps it generic rather than specific to Cedarline.
NO CREDIT, 0 TO 1	Claims it helps, with no explanation of how it works and no Cedarline outcome.

We're running AI workloads in cloud accounts, but we don't have a reliable way to see which models and AI assets are actually in use. We're also concerned about workloads reaching out to third-party AI models without security knowing about it. And even when we do know the traffic path, we need a practical way to inspect and enforce guardrails on prompt and response traffic in the cloud without changing every application.

Problem Identifier OUT OF 5

FULL, 5	Seeing AI assets and third-party model usage, and enforcing runtime guardrails on AI traffic without changing applications.
PARTIAL, 1 TO 4	Sees the concern but names only part of it, or states it in general terms without the specific requirement above.
NO CREDIT, 0	Restates or misreads the concern; the underlying requirement is missing.

Solution Designer OUT OF 10

FULL, 8 TO 10	Cisco AI Defense with Cisco Multicloud Defense (AI Runtime Protection) Also acceptable: Cisco Hypershield (runtime enforcement) (an HMF path to enforce guardrails on workload and AI traffic without changing applications); Cisco Multicloud Defense (cloud runtime enforcement) (enforces guardrails on AI traffic in the cloud alongside AI Defense).
PARTIAL, 2 TO 7	Names AI Defense or Multicloud Defense but not the runtime guardrail on prompt and response traffic.
NO CREDIT, 0 TO 1	Blocking AI use outright, or an answer with no runtime inspection or enforcement.

Value Communicator OUT OF 10

FULL, 8 TO 10	How it works. That's exactly what the Cisco AI Defense-Cisco Multicloud Defense integration is designed to address. Cisco AI Defense uses the integration to discover AI assets in your cloud environments and identify external AI destinations your workloads are calling. Outcome. Finds AI assets in cloud accounts. Identifies third-party AI model usage.
PARTIAL, 2 TO 7	Explains the how or the outcome, but not both, or keeps it generic rather than specific to Cedarline.
NO CREDIT, 0 TO 1	Claims it helps, with no explanation of how it works and no Cedarline outcome.

2

Case Study 2

For each pain point, grade all three dimensions. Each shows the **full-marks** answer, what a **partial** answer looks like, and what earns **no credit**, mapped to the score.

PAIN POINT 1 Grade all three dimensions, full marks 25

If a breach occurs in one of our public cloud instances, we are terrified that the attacker will move laterally into our core on-premises systems. How can we contain these threats in a multicloud environment?

Problem Identifier OUT OF 5

FULL, 5	Containing lateral movement between public cloud and on-prem in a multicloud environment.
PARTIAL, 1 TO 4	Sees the concern but names only part of it, or states it in general terms without the specific requirement above.
NO CREDIT, 0	Restates or misreads the concern; the underlying requirement is missing.

Solution Designer OUT OF 10

FULL, 8 TO 10	Cisco Multicloud Defense (macro and micro-segmentation gateways) Also acceptable: Cisco Secure Workload (micro-segmentation) (contains east-west movement at the workload; credit alongside or instead of the gateways when argued); Cisco Hypershield (distributed segmentation and enforcement across the estate).
PARTIAL, 2 TO 7	Names Multicloud Defense but not segmentation to contain lateral movement.
NO CREDIT, 0 TO 1	A single perimeter firewall with no multicloud or east-west control.

Value Communicator OUT OF 10

FULL, 8 TO 10	How it works. Cisco Multicloud Defense leverages macro and micro-segmentation capabilities to enforce strict communication boundaries. By integrating with our Hybrid Mesh Firewall and Hypershield, we can isolate workloads and inspect traffic between cloud segments, effectively stopping lateral movement before it reaches your sensitive data. Outcome. It provides granular visibility and control over east-west traffic. By applying Zero Trust principles across the hybrid fabric, you transform your network into a series of secure, isolated zones that prevent an attacker from pivoting through your infrastructure.
PARTIAL, 2 TO 7	Explains the how or the outcome, but not both, or keeps it generic rather than specific to Cedarline.
NO CREDIT, 0 TO 1	Claims it helps, with no explanation of how it works and no Cedarline outcome.

Our auditors require proof that our security controls are applied consistently across all our cloud workloads. Gathering this data from different cloud providers is a manual, time-consuming nightmare.

Problem Identifier OUT OF 5

FULL, 5	Consistent, auditable proof that security controls are applied across all cloud workloads.
PARTIAL, 1 TO 4	Sees the concern but names only part of it, or states it in general terms without the specific requirement above.
NO CREDIT, 0	Restates or misreads the concern; the underlying requirement is missing.

Solution Designer OUT OF 10

FULL, 8 TO 10	Cisco Multicloud Defense (unified logging and visibility) Also acceptable: Cisco Security Cloud Control (central visibility and reporting across the estate to evidence controls to auditors).
PARTIAL, 2 TO 7	Names Multicloud Defense but not the audit or evidence angle.
NO CREDIT, 0 TO 1	Per-cloud native logs or spreadsheets with no unified, auditable view.

Value Communicator OUT OF 10

FULL, 8 TO 10	How it works. Cisco Multicloud Defense provides unified logging and traceability. Because our security framework is integrated across the network fabric, you can generate consolidated reports that map directly to your compliance requirements (such as NIS2 or DORA) across all cloud environments from a single dashboard. Outcome. It automates the collection of evidence for security audits. By consolidating logs and policy enforcement data, you create a "single source of truth" that simplifies compliance reporting and demonstrates proactive governance to regulators.
PARTIAL, 2 TO 7	Explains the how or the outcome, but not both, or keeps it generic rather than specific to Cedarline.
NO CREDIT, 0 TO 1	Claims it helps, with no explanation of how it works and no Cedarline outcome.

We have workloads spread across AWS, Azure, and on-premises data centers. Managing security policies individually in each cloud environment is creating massive operational overhead and leading to inconsistent security posture.

Problem Identifier OUT OF 5

FULL, 5	One consistent policy model across AWS, Azure and on-prem, without per-cloud overhead.
PARTIAL, 1 TO 4	Sees the concern but names only part of it, or states it in general terms without the specific requirement above.
NO CREDIT, 0	Restates or misreads the concern; the underlying requirement is missing.

Solution Designer OUT OF 10

FULL, 8 TO 10	Cisco Multicloud Defense (unified policy across public clouds) Also acceptable: Cisco Secure Firewall with FMC or Security Cloud Control (consistent policy across the firewall estate; credit when the one-policy-model reasoning fits Cedarline).
PARTIAL, 2 TO 7	Names Multicloud Defense but not the single, consistent policy model across clouds.
NO CREDIT, 0 TO 1	Separate per-cloud native firewalls managed independently.

Value Communicator OUT OF 10

FULL, 8 TO 10	How it works. Cisco Multicloud Defense, integrated within our Hybrid Mesh Firewall architecture, provides a unified control plane. You can define security policies once via Security Cloud Control and enforce them consistently across your entire distributed environment, regardless of the underlying cloud provider. Outcome. It replaces siloed, manual configuration with centralized orchestration. By abstracting the security policy from the cloud-native tools, it ensures that your security posture is uniform and eliminates the "configuration drift" that occurs when managing multiple cloud consoles.
PARTIAL, 2 TO 7	Explains the how or the outcome, but not both, or keeps it generic rather than specific to Cedarline.
NO CREDIT, 0 TO 1	Claims it helps, with no explanation of how it works and no Cedarline outcome.

We have workloads on-prem, in cloud, and in mixed environments, and policy consistency is becoming difficult. Every environment feels different, and we're worried about drift.

Problem Identifier OUT OF 5

FULL, 5	Unified segmentation that prevents policy drift across hybrid and multicloud.
PARTIAL, 1 TO 4	Sees the concern but names only part of it, or states it in general terms without the specific requirement above.
NO CREDIT, 0	Restates or misreads the concern; the underlying requirement is missing.

Solution Designer OUT OF 10

FULL, 8 TO 10	Cisco Secure Workload with Cisco Secure Firewall (managed via FMC and Security Cloud Control) Also acceptable: Cisco Multicloud Defense (consistent segmentation policy) (enforces uniform segmentation across clouds; credit with Secure Workload's dependency mapping); Cisco Hypershield (distributed, consistent enforcement that prevents policy drift).
PARTIAL, 2 TO 7	Names Secure Workload or a firewall but not the drift-prevention or unified policy.
NO CREDIT, 0 TO 1	Point segmentation in one environment only.

Value Communicator OUT OF 10

FULL, 8 TO 10	How it works. That's a common hybrid environment challenge. Cisco Secure Workload helps create a more unified segmentation model by using application dependency awareness and AI/ML-driven policy discovery to define policy based on how workloads actually communicate. Outcome. Reduces policy drift across hybrid and multicloud environments. Aligns policy to application behavior, not just static IP-based rules.
PARTIAL, 2 TO 7	Explains the how or the outcome, but not both, or keeps it generic rather than specific to Cedarline.
NO CREDIT, 0 TO 1	Claims it helps, with no explanation of how it works and no Cedarline outcome.

3

Case Study 3

For each pain point, grade all three dimensions. Each shows the **full-marks** answer, what a **partial** answer looks like, and what earns **no credit**, mapped to the score.

PAIN POINT 1 Grade all three dimensions, full marks 25

When our security tools detect a compromised device, our manual response time is too slow. By the time we identify the user and manually isolate the port, the attacker has already moved through the network.

Problem Identifier OUT OF 5

FULL, 5	Automated, real-time isolation of a compromised device, faster than a manual response.
PARTIAL, 1 TO 4	Sees the concern but names only part of it, or states it in general terms without the specific requirement above.
NO CREDIT, 0	Restates or misreads the concern; the underlying requirement is missing.

Solution Designer OUT OF 10

FULL, 8 TO 10	Cisco ISE (Security Group Tags via pxGrid and Adaptive Network Control) Also acceptable: Cisco Secure Firewall integrated with ISE (firewall-triggered containment through ISE; credit when the automated quarantine path is clear); Cisco ISE with Adaptive Network Control (the core real-time auto-quarantine capability).
PARTIAL, 2 TO 7	Names ISE but not the automated, real-time quarantine through pxGrid or ANC.
NO CREDIT, 0 TO 1	A manual, ticket-based response, or a firewall rule with no identity or automation.

Value Communicator OUT OF 10

FULL, 8 TO 10	How it works. Cisco ISE integrates with your security ecosystem via pxGrid and Adaptive Network Control (ANC). When a threat is detected, ISE can automatically trigger a quarantine policy, isolating the compromised device in real-time without requiring manual intervention. Outcome. It automates the "Response" phase of the incident lifecycle. By bridging the gap between detection and enforcement, it stops threats at the speed of the network.
PARTIAL, 2 TO 7	Explains the how or the outcome, but not both, or keeps it generic rather than specific to Cedarline.
NO CREDIT, 0 TO 1	Claims it helps, with no explanation of how it works and no Cedarline outcome.

Our firewall environment has become too complex to manage manually.

Problem Identifier OUT OF 5

FULL, 5	Simplifying and cleaning up firewall policy administration at scale.
PARTIAL, 1 TO 4	Sees the concern but names only part of it, or states it in general terms without the specific requirement above.
NO CREDIT, 0	Restates or misreads the concern; the underlying requirement is missing.

Solution Designer OUT OF 10

FULL, 8 TO 10	Cisco Security Cloud Control with FMC and Cisco Secure Firewall (unified policy administration) Also acceptable: Cisco Secure Firewall Management Center (FMC) (centralised policy administration and rulebase cleanup; credit when the at-scale simplification is argued).
PARTIAL, 2 TO 7	Names a manager but not the rulebase cleanup or consolidation at scale.
NO CREDIT, 0 TO 1	Editing firewalls device by device with no central management.

Value Communicator OUT OF 10

FULL, 8 TO 10	How it works. Cisco applies AI-driven analysis to firewall telemetry and policy data to surface optimization opportunities, reduce manual troubleshooting, and guide cleaner policy administration. Outcome. Reduces manual policy review; identifies redundant/stale rules; lowers risk from human error.
PARTIAL, 2 TO 7	Explains the how or the outcome, but not both, or keeps it generic rather than specific to Cedarline.
NO CREDIT, 0 TO 1	Claims it helps, with no explanation of how it works and no Cedarline outcome.

Most of our traffic is encrypted, so we've lost visibility into a huge part of our environment.

Problem Identifier OUT OF 5

FULL, 5	Restoring visibility into encrypted traffic without decrypting everything.
PARTIAL, 1 TO 4	Sees the concern but names only part of it, or states it in general terms without the specific requirement above.
NO CREDIT, 0	Restates or misreads the concern; the underlying requirement is missing.

Solution Designer OUT OF 10

FULL, 8 TO 10	Cisco Secure Firewall with the Encrypted Visibility Engine (EVE) Also acceptable: Cisco Secure Firewall with selective TLS decryption (a valid but heavier path; credit when selective decryption is justified, though EVE is preferred for no-decrypt visibility).
PARTIAL, 2 TO 7	Names Secure Firewall but not EVE or a no-decrypt approach to encrypted traffic.
NO CREDIT, 0 TO 1	Blanket decryption of all traffic, or ignoring encrypted flows entirely.

Value Communicator OUT OF 10

FULL, 8 TO 10	How it works. Encrypted Visibility Engine (EVE) uses AI/ML-based fingerprinting to analyze encrypted traffic patterns and identify suspicious behavior without decrypting the payload. Outcome. Restores visibility into encrypted blind spots; avoids the performance/privacy overhead of decrypt-everything strategies.
PARTIAL, 2 TO 7	Explains the how or the outcome, but not both, or keeps it generic rather than specific to Cedarline.
NO CREDIT, 0 TO 1	Claims it helps, with no explanation of how it works and no Cedarline outcome.

We're worried traditional signature-based detection won't catch new or unknown attacks quickly enough.

Problem Identifier OUT OF 5

FULL, 5	Detecting unknown and zero-day threats beyond signature-based detection.
PARTIAL, 1 TO 4	Sees the concern but names only part of it, or states it in general terms without the specific requirement above.
NO CREDIT, 0	Restates or misreads the concern; the underlying requirement is missing.

Solution Designer OUT OF 10

FULL, 8 TO 10	Cisco Secure Firewall with Snort ML Also acceptable: Cisco Secure Firewall with Talos-backed IPS and malware protection (complementary advanced detection; credit when the beyond-signature reasoning is present, with Snort ML preferred).
PARTIAL, 2 TO 7	Names Secure Firewall or IPS but not machine-learning or zero-day detection.
NO CREDIT, 0 TO 1	Signature-only IPS with no unknown-threat detection.

Value Communicator OUT OF 10

FULL, 8 TO 10	How it works. That's exactly where Snort ML helps. It adds machine-learning-based detection inside the Snort 3 engine to identify exploit behavior associated with unknown/zero-day threats. Outcome. Addresses the gap between known and unknown threat detection; reduces dependence on signature timing.
PARTIAL, 2 TO 7	Explains the how or the outcome, but not both, or keeps it generic rather than specific to Cedarline.
NO CREDIT, 0 TO 1	Claims it helps, with no explanation of how it works and no Cedarline outcome.

4

Case Study 4

For each pain point, grade all three dimensions. Each shows the **full-marks** answer, what a **partial** answer looks like, and what earns **no credit**, mapped to the score.

PAIN POINT 1 Grade all three dimensions, full marks 25

Network policy helps, but we also need to know what's happening at runtime inside the workload. We need better detection of suspicious behavior and policy violations.

Problem Identifier OUT OF 5

FULL, 5	Runtime detection of suspicious behaviour and policy violations inside workloads.
PARTIAL, 1 TO 4	Sees the concern but names only part of it, or states it in general terms without the specific requirement above.
NO CREDIT, 0	Restates or misreads the concern; the underlying requirement is missing.

Solution Designer OUT OF 10

FULL, 8 TO 10	Isovalent Tetragon (eBPF runtime security) Also acceptable: Cisco Secure Workload (process and behaviour visibility with policy for workloads; credit alongside or instead of Tetragon when argued); Cisco Hypershield (in-kernel eBPF runtime detection and enforcement).
PARTIAL, 2 TO 7	Names an eBPF or workload tool but not the runtime behaviour detection inside workloads.
NO CREDIT, 0 TO 1	Perimeter intrusion prevention only, with no workload runtime visibility.

Value Communicator OUT OF 10

FULL, 8 TO 10	How it works. That's where Tetragon comes in. Tetragon extends the value of Cilium by providing eBPF-based runtime security and threat detection. Outcome. Adds runtime visibility beyond basic network controls. Detects suspicious system and network behavior in real time.
PARTIAL, 2 TO 7	Explains the how or the outcome, but not both, or keeps it generic rather than specific to Cedarline.
NO CREDIT, 0 TO 1	Claims it helps, with no explanation of how it works and no Cedarline outcome.

Our Kubernetes networking feels fragmented across clusters and environments. We're dealing with too many moving parts, inconsistent policies, and too much operational complexity.

Problem Identifier OUT OF 5

FULL, 5	Consistent Kubernetes networking and policy across clusters, with less operational complexity.
PARTIAL, 1 TO 4	Sees the concern but names only part of it, or states it in general terms without the specific requirement above.
NO CREDIT, 0	Restates or misreads the concern; the underlying requirement is missing.

Solution Designer OUT OF 10

FULL, 8 TO 10	Isovalent Cilium (eBPF Kubernetes networking and policy) Also acceptable: Cisco Hypershield (distributed eBPF enforcement across Kubernetes and workloads).
PARTIAL, 2 TO 7	Names Cilium but not the consistent networking and policy across clusters.
NO CREDIT, 0 TO 1	Per-cluster bespoke networking with no consistency.

Value Communicator OUT OF 10

FULL, 8 TO 10	How it works. That's exactly where Isovalent Cilium helps. Cilium provides a modern, eBPF-based Kubernetes networking layer that gives you consistent connectivity, policy enforcement, and visibility across clusters and environments. Outcome. Replaces brittle, IP-centric networking models with a Kubernetes-native approach. Improves consistency across dynamic cluster environments.
PARTIAL, 2 TO 7	Explains the how or the outcome, but not both, or keeps it generic rather than specific to Cedarline.
NO CREDIT, 0 TO 1	Claims it helps, with no explanation of how it works and no Cedarline outcome.

Right now we're stitching together separate tools for networking, security, and observability in Kubernetes, and it's creating too much overhead.

Problem Identifier OUT OF 5

FULL, 5	Consolidating Kubernetes networking, security and observability to reduce tool sprawl.
PARTIAL, 1 TO 4	Sees the concern but names only part of it, or states it in general terms without the specific requirement above.
NO CREDIT, 0	Restates or misreads the concern; the underlying requirement is missing.

Solution Designer OUT OF 10

FULL, 8 TO 10	Isovalent Cilium with Tetragon Also acceptable: Cisco Hypershield with Cisco Secure Workload (a consolidation path across enforcement and dependency visibility; credit when the tool-sprawl reduction is argued).
PARTIAL, 2 TO 7	Names Cilium or Tetragon alone but not the consolidation of networking, security and observability.
NO CREDIT, 0 TO 1	Adding another point tool that increases sprawl.

Value Communicator OUT OF 10

FULL, 8 TO 10	How it works. That's exactly the consolidation opportunity Isovalent addresses. With Cilium for networking and policy, and Tetragon for runtime security, you get a more unified platform for connectivity, segmentation, observability, and security. Outcome. Consolidates multiple functions into one platform approach. Reduces tool sprawl and operational friction.
PARTIAL, 2 TO 7	Explains the how or the outcome, but not both, or keeps it generic rather than specific to Cedarline.
NO CREDIT, 0 TO 1	Claims it helps, with no explanation of how it works and no Cedarline outcome.

We know we need microsegmentation, but we don't have enough visibility into east-west flows or application dependencies to create policy safely. We don't want to break the application.

Problem Identifier OUT OF 5

FULL, 5	Mapping east-west flows and application dependencies to segment safely without breaking applications.
PARTIAL, 1 TO 4	Sees the concern but names only part of it, or states it in general terms without the specific requirement above.
NO CREDIT, 0	Restates or misreads the concern; the underlying requirement is missing.

Solution Designer OUT OF 10

FULL, 8 TO 10	Cisco Secure Workload (application dependency mapping) Also acceptable: Isovalent Cilium with Hubble (flow observability) (maps east-west flows to inform safe segmentation; credit alongside or instead of Secure Workload when argued).
PARTIAL, 2 TO 7	Names Secure Workload but not the dependency mapping before enforcement.
NO CREDIT, 0 TO 1	Enforcing segmentation with no dependency discovery, which risks breaking applications.

Value Communicator OUT OF 10

FULL, 8 TO 10	How it works. That's one of the biggest reasons customers adopt Cisco Secure Workload. It gives you deep visibility into workload communication and application dependency mapping, so you can see how applications actually behave before enforcing policy. Outcome. Reveals actual application communication paths. Reduces the risk of segmentation causing outages.
PARTIAL, 2 TO 7	Explains the how or the outcome, but not both, or keeps it generic rather than specific to Cedarline.
NO CREDIT, 0 TO 1	Claims it helps, with no explanation of how it works and no Cedarline outcome.

A completed score sheet EXAMPLE

This is the Proctor evaluation page from the back of the Workbook, filled in for an example team, so you can see how to record the scores and feedback before you hand the result back.

Team Blue Team Members (up to 8) 8 present Date 12 May

CASE STUDY	SCORE / 100	FEEDBACK FOR THIS CASE STUDY
Case Study 1	66	Right products throughout and strong on AI Defense. Tighten the Article 15 requirement and state the concrete outcomes.
Case Study 2	78	Excellent multicloud segmentation and a clear audit story. One pain point was light on the single, unified policy model.
Case Study 3	72	Good automated containment with ISE. The encrypted-traffic answer needed EVE rather than blanket decryption.
Case Study 4	64	Solid Kubernetes segmentation. Dependency mapping before enforcement was implied rather than stated.
Overall total	280 / 400	

Overall grade Established (Advanced 300 to 400 · Established 200 to 299 · Foundational up to 199)

Topology diagram: SME review. The team's Hybrid Mesh Firewall diagram is submitted for Subject Matter Expert review and is not part of the score. Reviewed by SME: Yes ☒ No ☐

What the team did well

Consistently named the right Cisco products and tied them to Cedarline's estate. Customer-ready value stories in Case Studies 2 and 3.

Where they can learn and grow

Name the underlying requirement before reaching for the product, and always close with the concrete Cedarline outcome.

That is the whole flow. Grade every pain point on the rubric, total each case study, then fill in this page inside the team's Workbook exactly like the example above, and hand it back as their result and feedback.