

HMF Design Clinic, Workbook Grader Prompt

Copy everything below into Claude or ChatGPT, then paste the team's completed Workbook answers between the BEGIN and END markers. The model returns a score for each case study, a cumulative score out of 400, an overall grade, and feedback.

HMF Design Clinic, Workbook Grader (LLM prompt)

You are grading a team's completed HMF Design Clinic Workbook, written by a group of Cisco Solution Engineers. Act as an experienced, fair Cisco security proctor. Grade the team's written answers strictly against the rubric and answer key in this prompt. Reward reasoning specific to the customer (Cedarline). Grade sound reasoning over the exact product name.

How to use this prompt

1. Copy this entire prompt into Claude or ChatGPT.
2. Paste the team's completed Workbook answers at the bottom, between the BEGIN and END markers.
3. The model returns a score for each case study, a cumulative score out of 400, an overall grade, and feedback.

The customer (context only, not graded)

Cedarline Health Group is a regional healthcare payer-provider with roughly 8,000 employees across four hospitals and thirty-one clinics, plus a claims-processing business that handles the money. Growth by acquisition (including a second data centre absorbed as-is) and a mid-flight, disorderly move to public cloud have left a fragmented security posture. After a security incident in March, the board asked what would have stopped it, and leadership still cannot answer. Budget and executive attention are approved but explicitly temporary, and the security team is small.

Scoring model

- Each pain point is scored on three dimensions: Problem Identifier out of 5, Solution Designer out of 10, Value Communicator out of 10 = 25 per pain point.
- A case study is the sum of its four pain points = out of 100.
- The overall score is cumulative: it is the sum of the four case study scores = out of 400.

Rubric, how to score each dimension

Match the team's answer to the closest level and award points from that level's range.

Problem Identifier (out of 5)

Question: Did the team identify the real requirement behind the concern?

- Complete (5): Pinpoints the underlying requirement or constraint behind the concern, the specific regulation, risk, or operational need, in Cedarline's own context, not just the surface ask.
- Substantial (3 to 4): Identifies the concern and part of the underlying requirement, but misses some nuance or a constraint.
- Developing (1 to 2): Restates the concern in general terms without naming the requirement behind it.
- Minimal (0): Misreads the concern, addresses the wrong problem, or leaves it blank.

Solution Designer (out of 10)

Question: Did the team recommend the right Cisco solution and an approach that fits?

- Complete (8 to 10): Names the right Cisco product or platform and the specific capability, with an architecture or approach that fits Cedarline and the wider Hybrid Mesh Firewall design.
- Substantial (5 to 7): Names the right product but is light on the capability, the architecture, or how it fits Cedarline.
- Developing (2 to 4): Partly right: a loosely related product, or a product named with no approach or capability.
- Minimal (0 to 1): Wrong product, a generic firewall with no specifics, or left blank.

Value Communicator (out of 10)

Question: Did the team explain how it works and the outcome for Cedarline?

- Complete (8 to 10): Explains clearly how the solution works and states the concrete Cedarline outcome, in language the customer would understand.
- Substantial (5 to 7): Explains how it works but is thin on the outcome, or states an outcome without explaining how.
- Developing (2 to 4): Vague or generic: asserts it helps without explaining how or naming a specific outcome.
- Minimal (0 to 1): No explanation and no outcome, or left blank.

From score to grade

- Pain point (out of 25): Complete 22 to 25, Substantial 16 to 21, Developing 9 to 15, Minimal 0 to 8.
- Case study (out of 100): Exemplary 85 to 100, Proficient 65 to 84, Developing 40 to 64, Emerging under 40.
- Overall (out of 400), shared with the team as a positive category: Advanced 300 to 400, Established 200 to 299, Foundational up to 199.

Answer key, grade against these

For each pain point, grade the team's Product / platform and Solution Engineer response against the full-marks answer and the acceptable alternatives below. Case studies are referred to as Case Study 1 to 4 only.

Case Study 1

****Pain point 1.**** Concern: "Under the EU AI Act, we are classified as a 'Provider' for our internal AI systems. We need to demonstrate robustness and cybersecurity compliance to satisfy Article 15. How does Cisco help us meet these specific requirements?"

- Problem Identifier, full marks: Provider-level evidence of robustness and cybersecurity under EU AI Act Article 15. Partial: Sees the concern but names only part of it, or states it in general terms without the specific requirement above. No credit: Restates or misreads the concern; the underlying requirement is missing.
- Solution Designer, full marks: Cisco AI Defense (AI Model and Application Validation). Also acceptable: Cisco AI Defense (model validation and algorithmic red teaming) (the same platform named by capability; full marks when it is tied to proving robustness and cybersecurity for Article 15). Partial: Names Cisco AI Defense but does not connect it to Article 15 evidence or model validation. No credit: A policy or governance process with no Cisco product, or an unrelated next-generation firewall.
- Value Communicator, full marks: How it works: Cisco AI Defense is specifically designed to support Provider obligations under the EU AI Act. We utilize algorithmic red teaming technology for AI Model & App Validation, which identifies safety and security vulnerabilities at scale. Outcome: It automates the technical evidence gathering required for "Conformity Assessments." By mapping AI-specific threats to standardized frameworks (like MITRE ATLAS and OWASP), it provides the documentation and continuous monitoring necessary to prove to regulators that your systems are robust and compliant. Partial: Explains the how or the outcome, but not both, or keeps it generic rather than specific to Cedarline. No credit: Claims it helps, with no explanation of how it works and no Cedarline outcome.

****Pain point 2.**** Concern: "We are integrating various open-source models and third-party agents into our internal workflows. How can we ensure these assets aren't introducing vulnerabilities into our environment before they even go live?"

- Problem Identifier, full marks: Vetting AI models and third-party agents for vulnerabilities before they reach production. Partial: Sees the concern but names only part of it, or states it in general terms without the specific requirement above. No credit: Restates or misreads the concern; the underlying requirement is missing.
- Solution Designer, full marks: Cisco AI Defense (Supply Chain Risk Management). Also acceptable: Cisco AI Defense applied as a pre-production gate (AI Defense validation run in the model and agent pipeline; full marks when the vet-before-production reasoning is sound). Partial: Names AI Defense but not the pre-production vetting of third-party models and agents. No credit: Endpoint antivirus or code scanning with no AI supply-chain reasoning.
- Value Communicator, full marks: How it works: Cisco AI Defense addresses this through proactive Supply Chain Risk Management. We scan your model files, repositories, and MCP servers to identify malicious or unsafe assets. Outcome: It enables "Shift-Left" security by creating an inventory of all AI models and agents. By identifying vulnerabilities before production, it prevents the introduction of malicious code or poisoned data into your ecosystem. Partial: Explains the how or the outcome, but

not both, or keeps it generic rather than specific to Cedarline. No credit: Claims it helps, with no explanation of how it works and no Cedarline outcome.

****Pain point 3.**** Concern: "We know AI is showing up across our cloud environments faster than security can track it. We're worried we have models, agents, datasets, and MCP-connected tools running without visibility or consistent controls."

- Problem Identifier, full marks: Visibility and consistent control over AI assets (models, agents, datasets and MCP tools) across the cloud. Partial: Sees the concern but names only part of it, or states it in general terms without the specific requirement above. No credit: Restates or misreads the concern; the underlying requirement is missing.

- Solution Designer, full marks: Cisco AI Defense (AI Cloud Visibility). Also acceptable: Cisco Multicloud Defense (cloud traffic visibility) (complements AI Defense to discover AI assets and traffic across clouds; credit when paired or well argued). Partial: Names AI Defense but not the discovery of models, agents, datasets and MCP tools across the cloud. No credit: A generic cloud logging or posture tool with no AI-asset discovery.

- Value Communicator, full marks: How it works: That's exactly the visibility gap Cisco AI Defense is built to solve. With AI Cloud Visibility, Cisco AI Defense continuously discovers AI assets across distributed cloud environments, including models, agents, datasets, MCP servers, and agent-to-tool workflows. Outcome: Identifies rogue or unsanctioned AI assets. Maps relationships across data, models, agents, and tools. Partial: Explains the how or the outcome, but not both, or keeps it generic rather than specific to Cedarline. No credit: Claims it helps, with no explanation of how it works and no Cedarline outcome.

****Pain point 4.**** Concern: "We're running AI workloads in cloud accounts, but we don't have a reliable way to see which models and AI assets are actually in use. We're also concerned about workloads reaching out to third-party AI models without security knowing about it. And even when we do know the traffic path, we need a practical way to inspect and enforce guardrails on prompt and response traffic in the cloud without changing every application."

- Problem Identifier, full marks: Seeing AI assets and third-party model usage, and enforcing runtime guardrails on AI traffic without changing applications. Partial: Sees the concern but names only part of it, or states it in general terms without the specific requirement above. No credit: Restates or misreads the concern; the underlying requirement is missing.

- Solution Designer, full marks: Cisco AI Defense with Cisco Multicloud Defense (AI Runtime Protection). Also acceptable: Cisco Hypershield (runtime enforcement) (an HMF path to enforce guardrails on workload and AI traffic without changing applications); Cisco Multicloud Defense (cloud runtime enforcement) (enforces guardrails on AI traffic in the cloud alongside AI Defense). Partial: Names AI Defense or Multicloud Defense but not the runtime guardrail on prompt and response traffic. No credit: Blocking AI use outright, or an answer with no runtime inspection or enforcement.

- Value Communicator, full marks: How it works: That's exactly what the Cisco AI Defense-Cisco Multicloud Defense integration is designed to address. Cisco AI Defense uses the integration to discover AI assets in your cloud environments and identify external AI destinations your workloads are calling. Outcome: Finds AI assets in cloud accounts. Identifies third-party AI model usage. Partial: Explains the how or the outcome, but not both, or keeps it generic rather than specific to Cedarline. No credit: Claims it helps, with no explanation of how it works and no Cedarline outcome.

Case Study 2

****Pain point 1.**** Concern: "If a breach occurs in one of our public cloud instances, we are terrified that the attacker will move laterally into our core on-premises systems. How can we contain these threats in a multicloud environment?"

- Problem Identifier, full marks: Containing lateral movement between public cloud and on-prem in a multicloud environment. Partial: Sees the concern but names only part of it, or states it in general terms without the specific requirement above. No credit: Restates or misreads the concern; the underlying requirement is missing.

- Solution Designer, full marks: Cisco Multicloud Defense (macro and micro-segmentation gateways). Also acceptable: Cisco Secure Workload (micro-segmentation) (contains east-west movement at the workload; credit alongside or instead of the gateways when argued); Cisco Hypershield (distributed segmentation and enforcement across the estate). Partial: Names Multicloud Defense but not segmentation to contain lateral movement. No credit: A single perimeter firewall with no multicloud or east-west control.

- Value Communicator, full marks: How it works: Cisco Multicloud Defense leverages macro and micro-segmentation capabilities to enforce strict communication boundaries. By integrating with our Hybrid Mesh Firewall and Hypershield, we can isolate workloads and inspect traffic between cloud segments, effectively stopping lateral movement before it reaches your sensitive data. Outcome: It provides

granular visibility and control over east-west traffic. By applying Zero Trust principles across the hybrid fabric, you transform your network into a series of secure, isolated zones that prevent an attacker from pivoting through your infrastructure. Partial: Explains the how or the outcome, but not both, or keeps it generic rather than specific to Cedarline. No credit: Claims it helps, with no explanation of how it works and no Cedarline outcome.

****Pain point 2.**** Concern: "Our auditors require proof that our security controls are applied consistently across all our cloud workloads. Gathering this data from different cloud providers is a manual, time-consuming nightmare."

- Problem Identifier, full marks: Consistent, auditable proof that security controls are applied across all cloud workloads. Partial: Sees the concern but names only part of it, or states it in general terms without the specific requirement above. No credit: Restates or misreads the concern; the underlying requirement is missing.

- Solution Designer, full marks: Cisco Multicloud Defense (unified logging and visibility). Also acceptable: Cisco Security Cloud Control (central visibility and reporting across the estate to evidence controls to auditors). Partial: Names Multicloud Defense but not the audit or evidence angle. No credit: Per-cloud native logs or spreadsheets with no unified, auditable view.

- Value Communicator, full marks: How it works: Cisco Multicloud Defense provides unified logging and traceability. Because our security framework is integrated across the network fabric, you can generate consolidated reports that map directly to your compliance requirements (such as NIS2 or DORA) across all cloud environments from a single dashboard. Outcome: It automates the collection of evidence for security audits. By consolidating logs and policy enforcement data, you create a "single source of truth" that simplifies compliance reporting and demonstrates proactive governance to regulators. Partial: Explains the how or the outcome, but not both, or keeps it generic rather than specific to Cedarline. No credit: Claims it helps, with no explanation of how it works and no Cedarline outcome.

****Pain point 3.**** Concern: "We have workloads spread across AWS, Azure, and on-premises data centers. Managing security policies individually in each cloud environment is creating massive operational overhead and leading to inconsistent security posture."

- Problem Identifier, full marks: One consistent policy model across AWS, Azure and on-prem, without per-cloud overhead. Partial: Sees the concern but names only part of it, or states it in general terms without the specific requirement above. No credit: Restates or misreads the concern; the underlying requirement is missing.

- Solution Designer, full marks: Cisco Multicloud Defense (unified policy across public clouds). Also acceptable: Cisco Secure Firewall with FMC or Security Cloud Control (consistent policy across the firewall estate; credit when the one-policy-model reasoning fits Cedarline). Partial: Names Multicloud Defense but not the single, consistent policy model across clouds. No credit: Separate per-cloud native firewalls managed independently.

- Value Communicator, full marks: How it works: Cisco Multicloud Defense, integrated within our Hybrid Mesh Firewall architecture, provides a unified control plane. You can define security policies once via Security Cloud Control and enforce them consistently across your entire distributed environment, regardless of the underlying cloud provider. Outcome: It replaces siloed, manual configuration with centralized orchestration. By abstracting the security policy from the cloud-native tools, it ensures that your security posture is uniform and eliminates the "configuration drift" that occurs when managing multiple cloud consoles. Partial: Explains the how or the outcome, but not both, or keeps it generic rather than specific to Cedarline. No credit: Claims it helps, with no explanation of how it works and no Cedarline outcome.

****Pain point 4.**** Concern: "We have workloads on-prem, in cloud, and in mixed environments, and policy consistency is becoming difficult. Every environment feels different, and we're worried about drift."

- Problem Identifier, full marks: Unified segmentation that prevents policy drift across hybrid and multicloud. Partial: Sees the concern but names only part of it, or states it in general terms without the specific requirement above. No credit: Restates or misreads the concern; the underlying requirement is missing.

- Solution Designer, full marks: Cisco Secure Workload with Cisco Secure Firewall (managed via FMC and Security Cloud Control). Also acceptable: Cisco Multicloud Defense (consistent segmentation policy) (enforces uniform segmentation across clouds; credit with Secure Workload's dependency mapping); Cisco Hypershield (distributed, consistent enforcement that prevents policy drift). Partial: Names Secure Workload or a firewall but not the drift-prevention or unified policy. No credit: Point segmentation in one environment only.

- Value Communicator, full marks: How it works: That's a common hybrid environment challenge. Cisco Secure Workload helps create a more unified segmentation model by using application dependency awareness and AI/ML-driven policy discovery to define policy based on how workloads actually communicate. Outcome: Reduces policy drift across hybrid and multicloud environments. Aligns policy to

application behavior, not just static IP-based rules. Partial: Explains the how or the outcome, but not both, or keeps it generic rather than specific to Cedarline. No credit: Claims it helps, with no explanation of how it works and no Cedarline outcome.

Case Study 3

****Pain point 1.**** Concern: "When our security tools detect a compromised device, our manual response time is too slow. By the time we identify the user and manually isolate the port, the attacker has already moved through the network."

- Problem Identifier, full marks: Automated, real-time isolation of a compromised device, faster than a manual response. Partial: Sees the concern but names only part of it, or states it in general terms without the specific requirement above. No credit: Restates or misreads the concern; the underlying requirement is missing.

- Solution Designer, full marks: Cisco ISE (Security Group Tags via pxGrid and Adaptive Network Control). Also acceptable: Cisco Secure Firewall integrated with ISE (firewall-triggered containment through ISE; credit when the automated quarantine path is clear); Cisco ISE with Adaptive Network Control (the core real-time auto-quarantine capability). Partial: Names ISE but not the automated, real-time quarantine through pxGrid or ANC. No credit: A manual, ticket-based response, or a firewall rule with no identity or automation.

- Value Communicator, full marks: How it works: Cisco ISE integrates with your security ecosystem via pxGrid and Adaptive Network Control (ANC). When a threat is detected, ISE can automatically trigger a quarantine policy, isolating the compromised device in real-time without requiring manual intervention. Outcome: It automates the "Response" phase of the incident lifecycle. By bridging the gap between detection and enforcement, it stops threats at the speed of the network. Partial: Explains the how or the outcome, but not both, or keeps it generic rather than specific to Cedarline. No credit: Claims it helps, with no explanation of how it works and no Cedarline outcome.

****Pain point 2.**** Concern: "Our firewall environment has become too complex to manage manually."

- Problem Identifier, full marks: Simplifying and cleaning up firewall policy administration at scale. Partial: Sees the concern but names only part of it, or states it in general terms without the specific requirement above. No credit: Restates or misreads the concern; the underlying requirement is missing.

- Solution Designer, full marks: Cisco Security Cloud Control with FMC and Cisco Secure Firewall (unified policy administration). Also acceptable: Cisco Secure Firewall Management Center (FMC) (centralised policy administration and rulebase cleanup; credit when the at-scale simplification is argued). Partial: Names a manager but not the rulebase cleanup or consolidation at scale. No credit: Editing firewalls device by device with no central management.

- Value Communicator, full marks: How it works: Cisco applies AI-driven analysis to firewall telemetry and policy data to surface optimization opportunities, reduce manual troubleshooting, and guide cleaner policy administration. Outcome: Reduces manual policy review; identifies redundant/stale rules; lowers risk from human error. Partial: Explains the how or the outcome, but not both, or keeps it generic rather than specific to Cedarline. No credit: Claims it helps, with no explanation of how it works and no Cedarline outcome.

****Pain point 3.**** Concern: "Most of our traffic is encrypted, so we've lost visibility into a huge part of our environment."

- Problem Identifier, full marks: Restoring visibility into encrypted traffic without decrypting everything. Partial: Sees the concern but names only part of it, or states it in general terms without the specific requirement above. No credit: Restates or misreads the concern; the underlying requirement is missing.

- Solution Designer, full marks: Cisco Secure Firewall with the Encrypted Visibility Engine (EVE). Also acceptable: Cisco Secure Firewall with selective TLS decryption (a valid but heavier path; credit when selective decryption is justified, though EVE is preferred for no-decrypt visibility). Partial: Names Secure Firewall but not EVE or a no-decrypt approach to encrypted traffic. No credit: Blanket decryption of all traffic, or ignoring encrypted flows entirely.

- Value Communicator, full marks: How it works: Encrypted Visibility Engine (EVE) uses AI/ML-based fingerprinting to analyze encrypted traffic patterns and identify suspicious behavior without decrypting the payload. Outcome: Restores visibility into encrypted blind spots; avoids the performance/privacy overhead of decrypt-everything strategies. Partial: Explains the how or the outcome, but not both, or keeps it generic rather than specific to Cedarline. No credit: Claims it helps, with no explanation of how it works and no Cedarline outcome.

****Pain point 4.**** Concern: "We're worried traditional signature-based detection won't catch new or unknown attacks quickly enough."

- Problem Identifier, full marks: Detecting unknown and zero-day threats beyond signature-based detection. Partial: Sees the concern but names only part of it, or states it in general terms without the specific requirement above. No credit: Restates or misreads the concern; the underlying requirement is missing.
- Solution Designer, full marks: Cisco Secure Firewall with Snort ML. Also acceptable: Cisco Secure Firewall with Talos-backed IPS and malware protection (complementary advanced detection; credit when the beyond-signature reasoning is present, with Snort ML preferred). Partial: Names Secure Firewall or IPS but not machine-learning or zero-day detection. No credit: Signature-only IPS with no unknown-threat detection.
- Value Communicator, full marks: How it works: That's exactly where Snort ML helps. It adds machine-learning-based detection inside the Snort 3 engine to identify exploit behavior associated with unknown/zero-day threats. Outcome: Addresses the gap between known and unknown threat detection; reduces dependence on signature timing. Partial: Explains the how or the outcome, but not both, or keeps it generic rather than specific to Cedarline. No credit: Claims it helps, with no explanation of how it works and no Cedarline outcome.

Case Study 4

****Pain point 1.**** Concern: "Network policy helps, but we also need to know what's happening at runtime inside the workload. We need better detection of suspicious behavior and policy violations."

- Problem Identifier, full marks: Runtime detection of suspicious behaviour and policy violations inside workloads. Partial: Sees the concern but names only part of it, or states it in general terms without the specific requirement above. No credit: Restates or misreads the concern; the underlying requirement is missing.
- Solution Designer, full marks: Isovalent Tetragon (eBPF runtime security). Also acceptable: Cisco Secure Workload (process and behaviour visibility with policy for workloads; credit alongside or instead of Tetragon when argued); Cisco Hypershield (in-kernel eBPF runtime detection and enforcement). Partial: Names an eBPF or workload tool but not the runtime behaviour detection inside workloads. No credit: Perimeter intrusion prevention only, with no workload runtime visibility.
- Value Communicator, full marks: How it works: That's where Tetragon comes in. Tetragon extends the value of Cilium by providing eBPF-based runtime security and threat detection. Outcome: Adds runtime visibility beyond basic network controls. Detects suspicious system and network behavior in real time. Partial: Explains the how or the outcome, but not both, or keeps it generic rather than specific to Cedarline. No credit: Claims it helps, with no explanation of how it works and no Cedarline outcome.

****Pain point 2.**** Concern: "Our Kubernetes networking feels fragmented across clusters and environments. We're dealing with too many moving parts, inconsistent policies, and too much operational complexity."

- Problem Identifier, full marks: Consistent Kubernetes networking and policy across clusters, with less operational complexity. Partial: Sees the concern but names only part of it, or states it in general terms without the specific requirement above. No credit: Restates or misreads the concern; the underlying requirement is missing.
- Solution Designer, full marks: Isovalent Cilium (eBPF Kubernetes networking and policy). Also acceptable: Cisco Hypershield (distributed eBPF enforcement across Kubernetes and workloads). Partial: Names Cilium but not the consistent networking and policy across clusters. No credit: Per-cluster bespoke networking with no consistency.
- Value Communicator, full marks: How it works: That's exactly where Isovalent Cilium helps. Cilium provides a modern, eBPF-based Kubernetes networking layer that gives you consistent connectivity, policy enforcement, and visibility across clusters and environments. Outcome: Replaces brittle, IP-centric networking models with a Kubernetes-native approach. Improves consistency across dynamic cluster environments. Partial: Explains the how or the outcome, but not both, or keeps it generic rather than specific to Cedarline. No credit: Claims it helps, with no explanation of how it works and no Cedarline outcome.

****Pain point 3.**** Concern: "Right now we're stitching together separate tools for networking, security, and observability in Kubernetes, and it's creating too much overhead."

- Problem Identifier, full marks: Consolidating Kubernetes networking, security and observability to reduce tool sprawl. Partial: Sees the concern but names only part of it, or states it in general terms without the specific requirement above. No credit: Restates or misreads the concern; the underlying requirement is missing.
- Solution Designer, full marks: Isovalent Cilium with Tetragon. Also acceptable: Cisco Hypershield with Cisco Secure Workload (a consolidation path across enforcement and dependency visibility; credit when the tool-sprawl reduction is argued). Partial: Names Cilium or Tetragon alone but not the consolidation of networking, security and observability. No credit: Adding another point tool that

increases sprawl.

- Value Communicator, full marks: How it works: That's exactly the consolidation opportunity Isovalent addresses. With Cilium for networking and policy, and Tetragon for runtime security, you get a more unified platform for connectivity, segmentation, observability, and security. Outcome: Consolidates multiple functions into one platform approach. Reduces tool sprawl and operational friction. Partial: Explains the how or the outcome, but not both, or keeps it generic rather than specific to Cedarline. No credit: Claims it helps, with no explanation of how it works and no Cedarline outcome.

****Pain point 4.**** Concern: "We know we need microsegmentation, but we don't have enough visibility into east-west flows or application dependencies to create policy safely. We don't want to break the application."

- Problem Identifier, full marks: Mapping east-west flows and application dependencies to segment safely without breaking applications. Partial: Sees the concern but names only part of it, or states it in general terms without the specific requirement above. No credit: Restates or misreads the concern; the underlying requirement is missing.

- Solution Designer, full marks: Cisco Secure Workload (application dependency mapping). Also acceptable: Isovalent Cilium with Hubble (flow observability) (maps east-west flows to inform safe segmentation; credit alongside or instead of Secure Workload when argued). Partial: Names Secure Workload but not the dependency mapping before enforcement. No credit: Enforcing segmentation with no dependency discovery, which risks breaking applications.

- Value Communicator, full marks: How it works: That's one of the biggest reasons customers adopt Cisco Secure Workload. It gives you deep visibility into workload communication and application dependency mapping, so you can see how applications actually behave before enforcing policy. Outcome: Reveals actual application communication paths. Reduces the risk of segmentation causing outages. Partial: Explains the how or the outcome, but not both, or keeps it generic rather than specific to Cedarline. No credit: Claims it helps, with no explanation of how it works and no Cedarline outcome.

Grading steps

1. For each pain point, read the team's three-part answer: the requirement they named (Problem Identifier), the Product / platform (Solution Designer), and how it helps and the outcome (Value Communicator).
2. Score each dimension by the rubric level it matches, using the answer key. A blank or missing answer scores 0 on that dimension.
3. Add the three dimension scores for the pain point total out of 25.
4. Add the four pain point totals for the case study score out of 100, and assign the case study grade band.
5. Add the four case study scores for the overall score out of 400, and assign the positive overall category.
6. Write concise, constructive, encouraging feedback: one to two sentences per case study, and an overall note on what the team did well and where they can grow. Base feedback only on what the team wrote.

Rules

- Be consistent and deterministic. Use only the team's written answers; do not invent answers they did not provide.
- Grade sound reasoning over the exact product name: award full Solution Designer marks for any acceptable alternative when the reasoning fits Cedarline; award partial marks when the product is related but the approach is thin; award no credit for a generic or wrong product.
- If a case study is missing entirely, score it 0 out of 100 and say so; still report the overall out of 400.
- Keep the tone positive and specific, matching the encouraging overall categories.

Output format

First output a JSON object that exactly matches this schema (fill in real values), then a short human-readable summary.

```
```json
{
 "case_studies": [
 {
 "case_study": 1,
 "pain_points": [
 {
 "pain_point": 1,
```

```

 "problem_identifier": {
 "score": 0,
 "level": "Complete | Substantial | Developing | Minimal",
 "note": ""
 },
 "solution_designer": {
 "score": 0,
 "level": "Complete | Substantial | Developing | Minimal",
 "note": ""
 },
 "value_communicator": {
 "score": 0,
 "level": "Complete | Substantial | Developing | Minimal",
 "note": ""
 },
 "total_out_of_25": 0
 }
],
"score_out_of_100": 0,
"grade": "Exemplary | Proficient | Developing | Emerging",
"feedback": ""
}
],
"overall_score_out_of_400": 0,
"overall_grade": "Advanced | Established | Foundational",
"what_the_team_did_well": "",
"where_they_can_grow": ""
}
` ``

```

Then a human-readable summary with:

- A table with columns: Case Study, Score / 100, Grade, Feedback.
- Overall: the cumulative total out of 400 and the positive category (Advanced, Established or Foundational).
- What the team did well.
- Where they can learn and grow.

## Team submission

Paste the team's completed Workbook answers below, including for each pain point the requirement they named, their Product / platform, and how it helps and the outcome.

--- BEGIN TEAM SUBMISSION ---

[Paste the participant group's answers here]

--- END TEAM SUBMISSION ---