

Grading Guide

Qualitative feedback with the Good · Better · Best ladder

A coaching tool, not a scoresheet. For each pain point you'll mark where the answer lands on the ladder and jot two things: what landed, and one nudge to level up. No numbers — just honest, encouraging, useful feedback.

Coach the answer up the ladder

Good

NAMES THE FIT

You put the right Cisco solution on the table and show you understand what the customer is actually worried about.

Better

EXPLAINS THE HOW

You connect that solution to the concern — how it works and why it answers this specific problem — in language the customer would nod along to.

Best

SELLS THE OUTCOME

You tie it to Cedarline's real constraints and the business outcome they care about, and you place it inside the bigger Hybrid Mesh Firewall story.

- 1 Read the team's **product** and **SE response** together for each pain point.
- 2 Mark where it lands — **Good**, **Better**, or **Best**. Use the **Solution Guide** for the model answer and the "what each level sounds like" cues.
- 3 Write **what landed** (always — every level earns recognition) and **one nudge** to reach the next rung.
- 4 At the end of each case study, capture a one-line **qualitative profile** and the team's strengths + growth area.

Feedback principles. Lead with strength — name what the team did well before what's missing. Be specific to Cedarline, not generic. One growth nudge beats five — pick the highest-leverage one. Meet each SE where they are: a generalist reaching **Good** is a genuine win.

Coaching both audiences. Specialists: push past the right product to the customer's specific constraints and the outcome. Generalists: affirm the fit they found, then hand them the one sentence that makes it **Better**.

1

CASE STUDY 1 · GRADING

AI Security

Mark each pain point on the ladder and note what landed + one nudge. The **Solution Guide** has the model answer and the Good/Better/Best cues for each concern.

PAIN POINT 1

Under the EU AI Act, we are classified as a 'Provider' for our internal AI systems. We need to demonstrate robustness and cybersecurity compliance to satisfy Article 15. How does Cisco help us meet these specific requirements?

Where it lands: ☐ Good ☐ Better ☐ Best

✓ What landed

↗ Level up

PAIN POINT 2

We are integrating various open-source models and third-party agents into our internal workflows. How can we ensure these assets aren't introducing vulnerabilities into our environment before they even go live?

Where it lands: ☐ Good ☐ Better ☐ Best

✓ What landed

↗ Level up

PAIN POINT 3

We know AI is showing up across our cloud environments faster than security can track it. We're worried we have models, agents, datasets, and MCP-connected tools running without visibility or consistent controls.

Where it lands: ☐ Good ☐ Better ☐ Best

✓ What landed

↗ Level up

PAIN POINT 4

We're running AI workloads in cloud accounts, but we don't have a reliable way to see which models and AI assets are actually in use. We're also concerned about workloads reaching out to third-party AI models without security knowing about it. And even when we do know the traffic path, we need a practical way to inspect and enforce guardrails on prompt and response traffic in the cloud without changing every application.

Where it lands: ☐ Good ☐ Better ☐ Best

✓ What landed

↗ Level up

Case Study 1 — qualitative profile

Overall this team is... (e.g. "mostly Better, reaching Best on segmentation")

BIGGEST STRENGTH

GROWTH AREA

2

CASE STUDY 2 · GRADING

Cloud Edge

Mark each pain point on the ladder and note what landed + one nudge. The **Solution Guide** has the model answer and the Good/Better/Best cues for each concern.

PAIN POINT 1

If a breach occurs in one of our public cloud instances, we are terrified that the attacker will move laterally into our core on-premises systems. How can we contain these threats in a multicloud environment?

Where it lands: ☐ Good ☐ Better ☐ Best

✓ What landed

↗ Level up

PAIN POINT 2

Our auditors require proof that our security controls are applied consistently across all our cloud workloads. Gathering this data from different cloud providers is a manual, time-consuming nightmare.

Where it lands: ☐ Good ☐ Better ☐ Best

✓ What landed

↗ Level up

PAIN POINT 3

We have workloads spread across AWS, Azure, and on-premises data centers. Managing security policies individually in each cloud environment is creating massive operational overhead and leading to inconsistent security posture.

Where it lands: ☐ Good ☐ Better ☐ Best

✓ What landed

↗ Level up

We have workloads on-prem, in cloud, and in mixed environments, and policy consistency is becoming difficult. Every environment feels different, and we're worried about drift.

Where it lands: ☐ Good ☐ Better ☐ Best

✓ What landed

↗ Level up

Case Study 2 — qualitative profile

Overall this team is... (e.g. "mostly Better, reaching Best on segmentation")

BIGGEST STRENGTH

GROWTH AREA

3

CASE STUDY 3 · GRADING

DC Edge · Perimeter Firewall

Mark each pain point on the ladder and note what landed + one nudge. The **Solution Guide** has the model answer and the Good/Better/Best cues for each concern.

PAIN POINT 1

When our security tools detect a compromised device, our manual response time is too slow. By the time we identify the user and manually isolate the port, the attacker has already moved through the network.

Where it lands: ☐ Good ☐ Better ☐ Best

✓ What landed

↗ Level up

PAIN POINT 2

Our firewall environment has become too complex to manage manually.

Where it lands: ☐ Good ☐ Better ☐ Best

✓ What landed

↗ Level up

PAIN POINT 3

Most of our traffic is encrypted, so we've lost visibility into a huge part of our environment.

Where it lands: ☐ Good ☐ Better ☐ Best

✓ What landed

↗ Level up

PAIN POINT 4

We're worried traditional signature-based detection won't catch new or unknown attacks quickly enough.

Where it lands: ☐ Good ☐ Better ☐ Best

✓ What landed

↗ Level up

Case Study 3 — qualitative profile

Overall this team is... (e.g. "mostly Better, reaching Best on segmentation")

BIGGEST STRENGTH

GROWTH AREA

4

CASE STUDY 4 • GRADING

Macro & Micro Segmentation

Mark each pain point on the ladder and note what landed + one nudge. The **Solution Guide** has the model answer and the Good/Better/Best cues for each concern.

PAIN POINT 1

Network policy helps, but we also need to know what's happening at runtime inside the workload. We need better detection of suspicious behavior and policy violations.

Where it lands: ☐ Good ☐ Better ☐ Best

✓ What landed

↗ Level up

PAIN POINT 2

Our Kubernetes networking feels fragmented across clusters and environments. We're dealing with too many moving parts, inconsistent policies, and too much operational complexity.

Where it lands: ☐ Good ☐ Better ☐ Best

✓ What landed

↗ Level up

PAIN POINT 3

Right now we're stitching together separate tools for networking, security, and observability in Kubernetes, and it's creating too much overhead.

Where it lands: ☐ Good ☐ Better ☐ Best

✓ What landed

↗ Level up

We know we need microsegmentation, but we don't have enough visibility into east-west flows or application dependencies to create policy safely. We don't want to break the application.

Where it lands: ☐ Good ☐ Better ☐ Best

✓ What landed

↗ Level up

Case Study 4 — qualitative profile

Overall this team is... (e.g. "mostly Better, reaching Best on segmentation")

BIGGEST STRENGTH

GROWTH AREA

The team's story across the clinic

Team _____ Event _____

✓ Strengths across the four case studies

↗ Priorities to develop

Overall qualitative profile

In a sentence — where is this team on the Good → Better → Best journey, and what's the one thing that would move them up?

Remember. The ladder describes the answer, not the person. Name the level, celebrate the win, and give one clear step up. Every team leaves knowing what they did well and exactly how to get better.